
Evolution: From Distributed Cache to Feature Store

Yao Yue, IOP Systems

About me

- 10+ years of working on cache, creator of [Pelikan](#)
- 5+ years of performance engineering
- IOP Systems, ex-Twitter (not to be confused with X/Twitter)
- Systems thinking enthusiast



Distributed Caching

- The emergence of distributed cache
 - Why
 - What
- Production adoption – distributed system grinding
 - Sharding and scaling
 - Availability
 - Performance
- Functionality
 - Data types and APIs

Fig. 1: Cache engine chronicles,
most likely partial

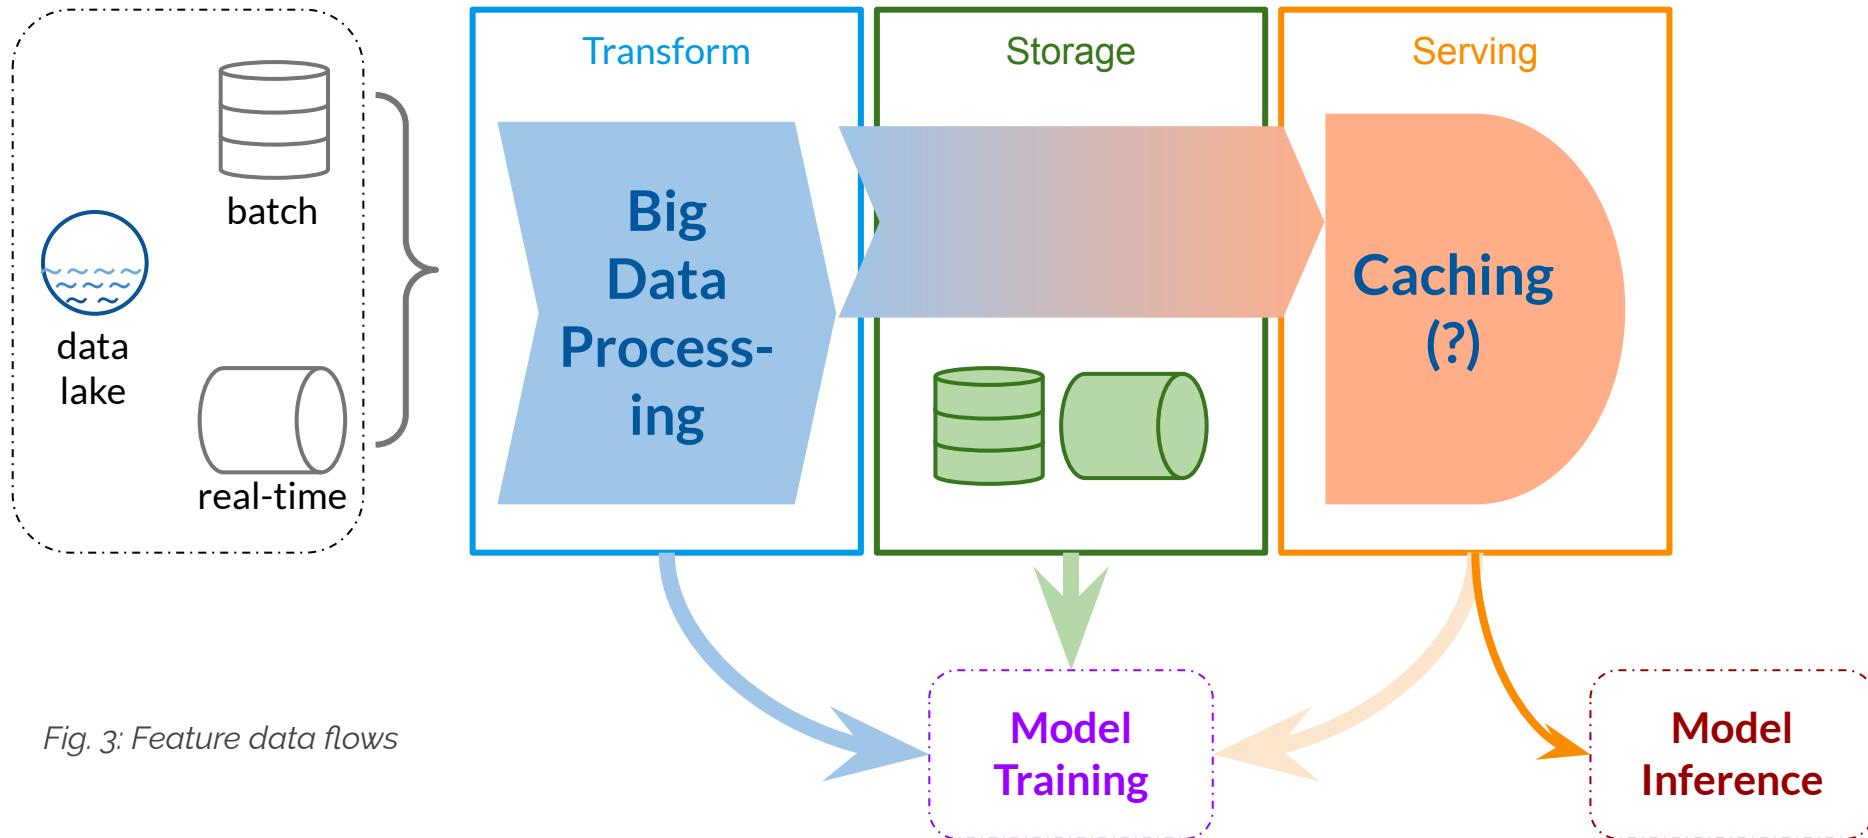


Big Data Processing

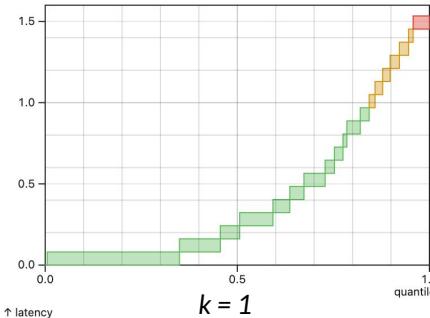
- First there was data, a lot of data
- Transform data with a distributed system
- An increasingly-online world requires streaming
- Unification of the batch / streaming processing
 - Compute semantics unification: Lambda architecture
 - Access semantics unification: Kappa architecture

Fig. 2: Data processing chronicle,
most definitely partial

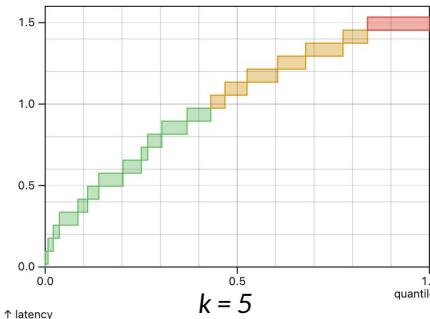
Feature store ecosystem



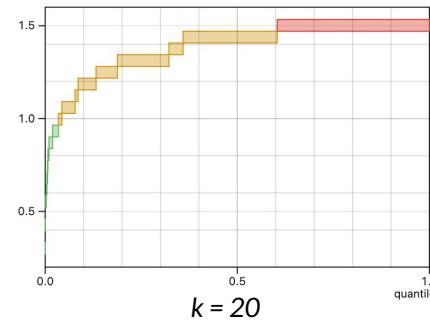
Bite-size big-data serving



$k = 1$



$k = 5$



$k = 20$

Identity-based feature retrieval

- A lot of data in a key-value format
- Not a canonical source
- Scatter-gather retrieval
 - Tail at scale: elevation of tail latencies

⇒ *Distributed caching*

➤ Data Layout

➤ Indexing

➤ API interface

➤ Threading
Architecture

Divergent read-write patterns

- Write
 - Batch / Streaming writes, from data processing jobs
 - vs *point update* in transactional workloads
 - High volume, sequential, delay-tolerant
- Read
 - Point reads
 - Low volume, random, latency-sensitive
 - (Streaming reads in the future?)

➤ Entry Encoding

➤ Native operators

➤ Programmability

Feature-specific Considerations

- Limited basic types
 - Floating-point, fixed-point, integers
- Organizations
 - Maps/Records
 - Sorted lists/maps
- Queries
 - `get(Multi)Field`
 - `range`
 - `op()` (`max`, `min`, `avg`, ...)

A few predictions on feature store evolution

- On storage
 - Embrace SSDs and write-friendly data layout, mainly for cost
 - Yes to durability, no to transactions
- On architecture
 - Local-level compute-storage separation
 - Message-passing centric internal communication
- Data import APIs
- High-level grammar for queries and functions

Questions?