

# Effective Performance Engineering at Twitter-Scale

IOP @ Twitter (2017.4 - 2022.11)  
Yao Yue, IOP Systems

# Performance Engineering...

---

“On the second night, he asked for a ladder, climbed up the generator and made a chalk mark on its side.”

“They did, and the generator performed to perfection.”

Charles Steinmetz, aka “Wizard of Schenectady”

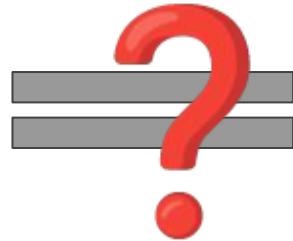


*(Steinmetz, itemizing the bill at the request of Henry Ford)*

Making chalk mark on generator \$1.

Knowing where to make mark \$9,999.

# Performance Engineering



# Wizardry

# Who I am

---

- ❑ Maintained & developed multiple distributed caching services
  - ❑ Trained by a decade of cache incidents, years of operations, system tuning, and capacity planning
  - ❑ Founded the performance team (IOP) at Twitter in 2017
- 
- ❑ Enjoy musing about (distributed) systems, and integration between emerging hardware and software trends

# Why we need performance engineering (more than ever)

For 30 years, hardware vendors had kept most performance engineers out of a job.

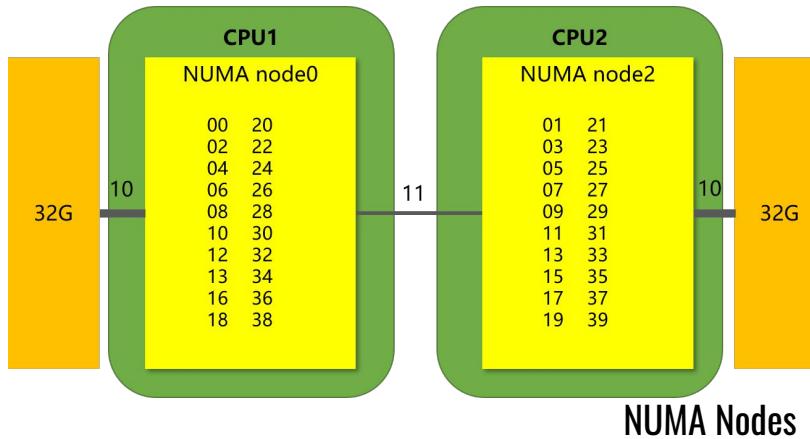


Now, our moment has finally come....

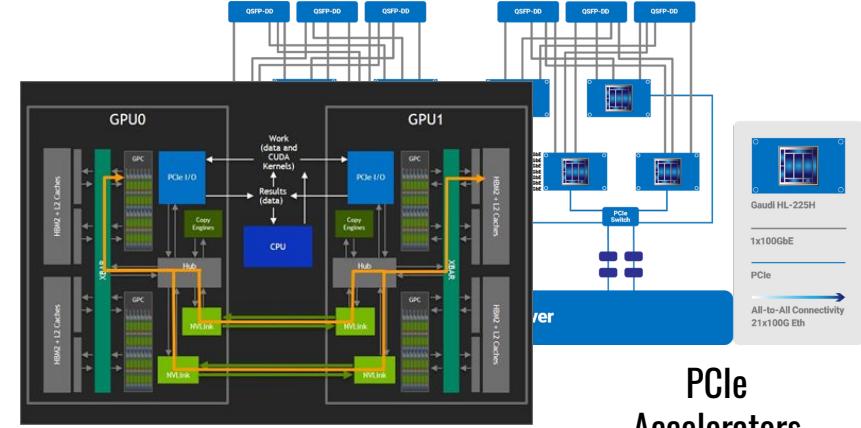
Power Cap

Thermal Cap

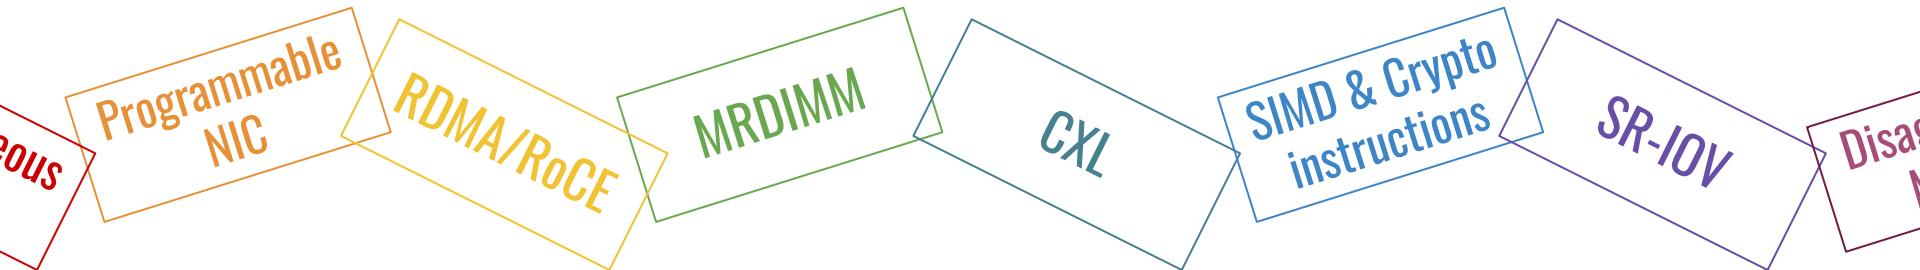
# Hardware Challenges — lunch is served, but hardly free



Credit: [github.com/LyleLee](https://github.com/LyleLee)



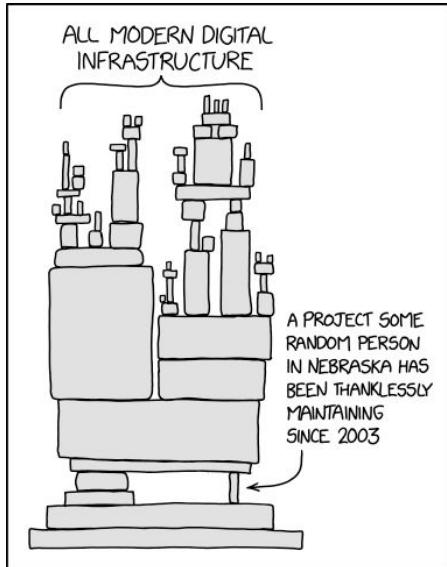
PCIe  
Accelerators



# Software Challenges — modern software is anti-hero

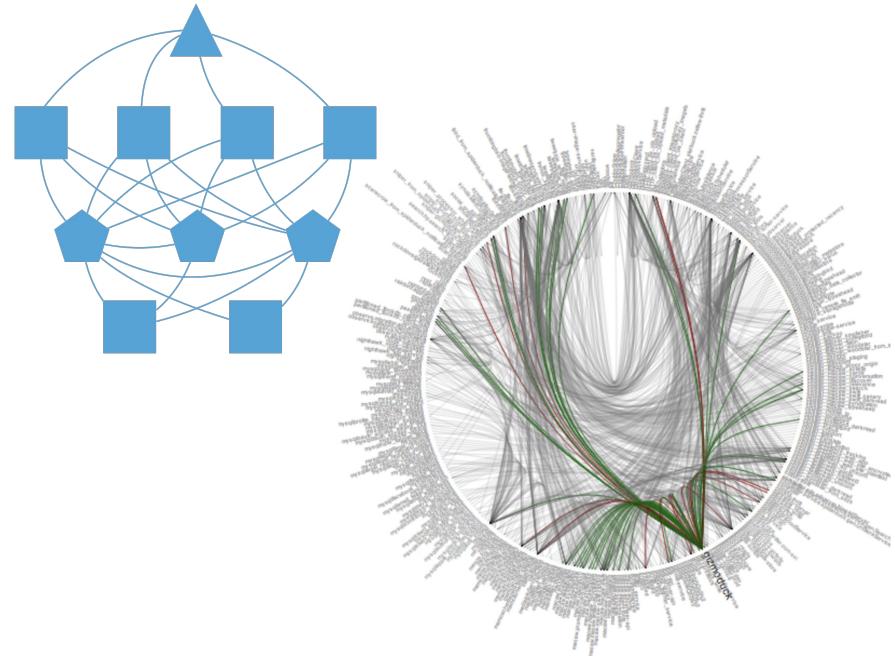
---

The stack is incredibly deep and wide.



Credit:  
[Randall Monroe](#)

The service ecosystem is a jungle – heterogeneous, diverse, and connected.



*It's the complexity,  
duh!*

# Performance is a System Property/Problem

---

**Systems** – models of (many) parts and relationships

**Performance** – a *counting* exercise

- of resources in a system
- at the right granularity
- that changes over time

Security	and
Availability	or
Performance	sum

System Property Intuitions

$$\text{Reliability} = \text{Availability} \times \text{Performance}$$

$$\text{Cost} = \text{Availability} \div \text{Performance}$$

# “How Slow Should It Be?”

---

- ❑ Priorities

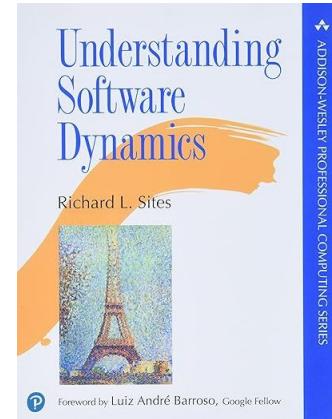
- Structure > Plumbing > Finish

- ❑ Predictability

- tail latency and consistency

*High performance services are all alike;*

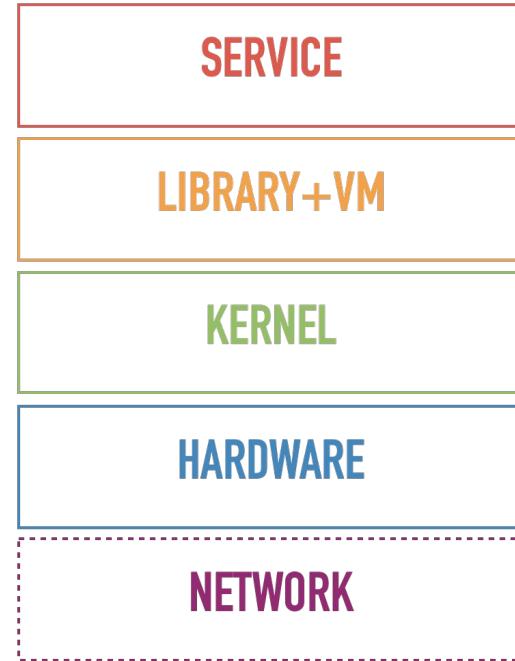
*every under-performing service is under-performing in its own way.*



# Performance Engineering at Scale

# Software Runtime as a Local System

- ❑ Characterize relationship between layers and components (“**System**”)
- ❑ Metrics for different resources and components (“**counting**”)
  - CPU, memory, network, storage...
  - syscall, scheduler...
  - GC
  - func call, ...



# How We Curated Data

---

- ❑ Rezolus: a performance telemetry agent

- Pluggable, largely eBPF-based samplers
- High frequency observability
- Extremely lightweight

Generating  
Signal

- ❑ Long Term Metrics (LTM): full-stack metrics

- From topline to hardware utilization
- Hardware, OS, container, JVM, Finagle, application, Rezolus, env metadata
- Minutely metric resolution with multi-year retention
- Normalization and verification
- All datasets query-ready with SQL

Aggregating  
Signal

# How We Used Data (to Get Insights)

---

- EasyPerf & EasyPerf UI for accessible GC wins
- Fleet health reports (a new team was born)
- Ranked utilization reports

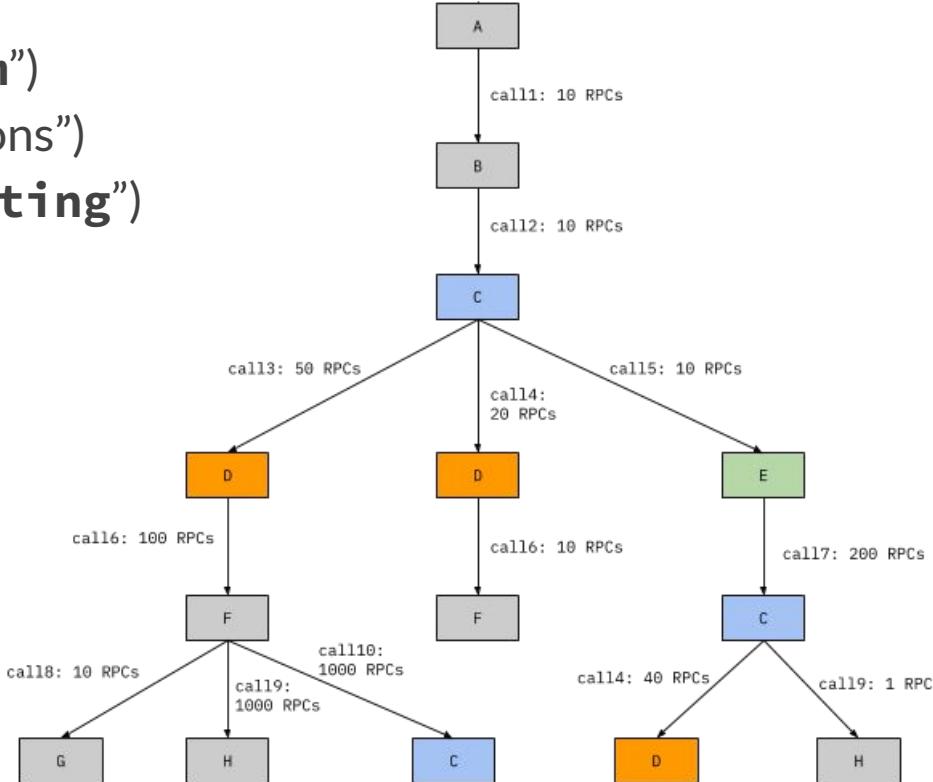
What the  
team did

- Capacity engineers derived utilization from LTM
- Service owners conducted performance analysis and added metrics
- HWEng, datacenter planners, and SiteOps preferred our machine metadata dataset over the canonical databases

What others  
did

# Application as a Distributed System

- ❑ Edges between services (“**System**”)
- ❑ Metrics and attributes (“annotations”) associated with each edge (“**counting**”)
  - send/recv timestamps
  - request/response sizes
  - status labels
  - (any free formed data)



# How We Curated Data

---

## ❑ Higher quality traces

- Tracing data is fraught with incomplete data, clock drift, ill-formatted fields...
- Continuous data validation and fixes
- Methodically set & recorded sampling rate
- Redesigned and re-implemented trace collection service

Generating  
Signal

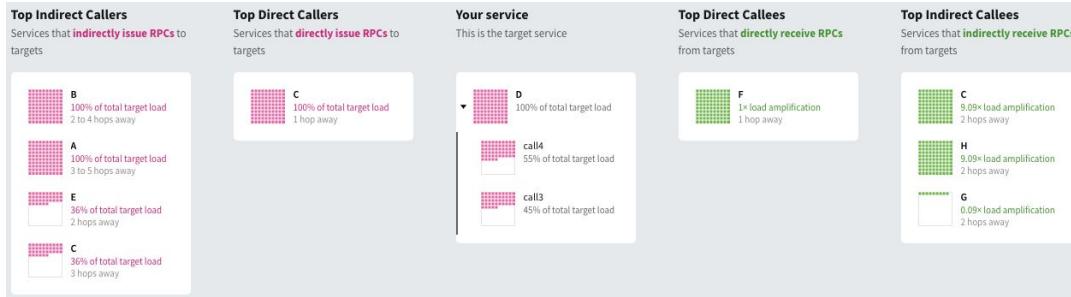
## ❑ Created a trace aggregation pipeline

- Separation of annotations and the call graph
- Multiple levels of indices: traces, span, annotation; and mappings between them
- All datasets query-ready with SQL

Aggregating  
Signal

# How We Used Data (to Get Insights)

- Service Dependency (Graph) Explorer



- Latenseer: a causal model of e2e latency distribution
- Many analyses on demand
- Data Privacy Team created a data propagation graph on top of service dependency graph

What the team did

What others did

# Impact of Our Work

---

- ❑ Eliminated the noisy neighbor problem via kernel changes
  - ❑ Better configuration & bin-packing of largest services
  - ❑ Direct optimization of largest services
  - ❑ Better hardware selection for largest services
  - ❑ >100MM savings overall
- 
- ❑ Datacenter and cloud migration readiness
  - ❑ Graceful degradation readiness
  - ❑ Numerous other analyses on-demand

# On Deck / In Progress

---

## Platforms

- ❑ Augur: Fleet-wide, continuous perf profiling and aggregation
- ❑ Systems Lab: a performance testing platform

## Products

- ❑ Trace Latency Visualizer

# How do Performance Engineers Fit In?

*They don't.*

# Questions Facing A Performance Engineering Team

---

## The Fundamentals

- Business value alignment
  - Topline
  - Bottomline
- Decision making structure
  - Top-down
  - Bottom-up

## Strategic Questions

- Who are my customers?
- Who is my champion?
- How to provide incentives?
- How to convince others?
- How to get recognized / promoted?

# Other Thoughts

---

- ❑ Embodying and fostering a culture of performance
  - ❑ Amplify and facilitate other people's performance work
  - ❑ Be a beacon, an advocate, and a warehouse of ideas
- ❑ The unique boom-and-bust cycle of performance engineering
  - ❑ Often trending the opposite way as the business
  - ❑ Solution: Trading between front- and back-stage work

# How it actually happened

There was *no*  
blueprint in the  
beginning.

*I was thinking  
~~really hard~~  
occasionally*

# The Philosopher's Approach to Performance

---

Does **performance** have meaning?

Can **performance** be judged objectively?

Is **performance** the same everywhere, all the time?

How can I find **performance**?

Am I alone in my pursuit for **performance**?

How do I share **performance**?



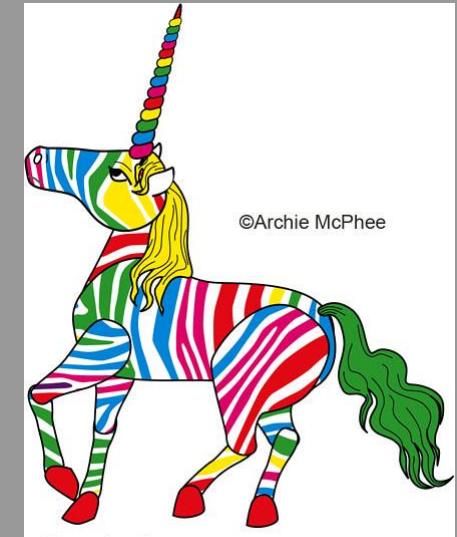
(ideal)



(reality)

# The ideal performance engineer —

has background in hardware,  
is good at reading/writing software,  
has dealt with large-scale operations,  
speaks fluent analytics,  
understands the language of business,  
can communicate effectively with all kinds  
of job roles and all levels of leadership. (Walks on water...)



We need a team.

We need a diverse team.

Find those that are  
[much] better than  
you [at some of the  
things].



# The First Phase (2017-2018, bootstrapping)

---

- ❑ Perf telemetry work commenced
- ❑ Lots of “odd jobs”
  - ❑ Consultation and debugging
  - ❑ Prototypes of ideas
  - ❑ Small wins
  - ❑ Favors (I somehow got involved in GDPR...)
- ❑ Writing down our vision & methodology

4 people  
All internal,  
SRE heavy

# The Second Phase (2018-19, laying foundations)

---

- ❑ Tracing work commenced
- ❑ Long Term Metrics work commenced
- ❑ Surveying new technologies for optimization opportunities
- ❑ Rezolus achieved initial success
- ❑ Offering more consultation

~8 people

Added hardware  
and analytics  
expertise

# The Third Phase (2020-2021, maturing)

---

- ❑ Primary datasets become usable
- ❑ Consulting now largely driven by inbound requests
- ❑ Frequently participated in key projects or crisis response
- ❑ Started to build more products
- ❑ Internal/External branding effort via talks and publishing

10+ people

Added UX and domain expertise

# The Fourth Phase (2022, an expansion cut abruptly)

---

- ❑ Coordinating several multi-team efficiency projects
- ❑ Planning and building the next batch of platform investments
- ❑ More investment into accessibility of our data via product and publications
- ❑ Negotiated an embedding model for localized performance engineering

~10 people  
Teams sprouted around us

# Lessons Learned

The technical and  
**social** considerations  
of performance  
engineering are  
equally important.

# Scalable Performance Methodology

---

- ❑ Strength in number: perf needs a signal-to-insight pipeline, and an opportunity-to-impact funnel
- ❑ Best places to look: infrastructure with broad impact, or areas with significant footprint
- ❑ Create platforms & products that enable a large number of individuals and teams to do similar work on their own
- ❑ Design the team to fit the broader organizational
- ❑ Outreach is serious work: education, consulting, collaboration, embedding, credit sharing...

People make work  
happen. Each person's  
strengths and  
personality matter.

# Building a Good Performance Team

---

- Treasure excellence
- Seek diversity
- Embrace chance
- Be helpful, be generous, make friends, build a culture

What happened to IOP after  
Twitter?

# IOP SYSTEMS



Yao Yue



Mihir Nanavati



Yuri Vishnevsky



Xi Yang



Brian Martin



Sean Lynch

Turning performance engineering expertise into platforms and products that could benefit everyone.

# Let's talk about Performance

---

**Yao Yue**

[yao@iop.systems](mailto:yao@iop.systems)

<https://iop.systems>



# Appendix

# Who We Were

---

**Full-timers:** Yao Yue, Brian Martin, Anatole Shaw, Andy Wilcox, Rebecca Isaacs, Dan Luu, Jonathan Simms, Kunal Trivedi, Michael Leinartas, Xi Yang, Mihir Nanavati, Sean Lynch

**Part-timers:** Yuri Vishnevsky, Oleksandr Kuvshynov

**Interns:** Juncheng Yang (x2), Sean Lynch (x2), Danieil Skrinikov, Lexiang Huang, Anna Karanika, Yazhuo Zhang

And many, many friends throughout Twitter 1.0: SRE, KaOS, HWEng, CSL, Capacity Planning, Revenue, Fleet Health, Kite (Service Identity & Chargeback), SiteOps, Core Services, TPM, Finance, ...

# What We Published

---

## Open Source

[Rezolus](#), [rpc-perf](#), [Pelikan Cache](#), [cache-trace](#), [Density Plot](#), misc [Rust utilities](#), ...

## Blogs

[Metric analytics](#), [Trace analytics](#), [Cache incidents](#), [Using ADQ](#), [Rezolus](#)

## Papers

[OSDI'20 \(TOS V17.3\)](#), [NSDI'21](#), [FAST'23](#), [HotOS'23](#), [SOSP'23](#), [SoCC'23](#), [NSDI'24](#)

## Talks

[YOW! 2017](#), [SNIA PMEM Summit 2020](#), [SDC 2020 \(keynote\)](#), [Strange Loop 2021](#), [P99 Conf 2022](#), [QConSF 2022](#)