

Why I left Big-*O* in the dustbin when I left school

Yao Yue

Why was Big-*O* introduced into computer science?

“To approximate and classify algorithms according to how their run time or space requirements grow as the input size grows”

-wikipedia

Is it true?

Binary search is *faster* than linear search: $O(\log n) < O(n)$

Hashmap insertion is *faster* than skip list: $O(1) < O(\log n)$

Coppersmith-Winograd is *faster* than naive matrix multiplication: $O(n^{2.37}) < O(n^3)$

As a member of the
Bachmann–Landau
notations, $O()$
predates computers.

The big lie:

scaling is smooth

Why Big-*O* fails to predict reality

Big-*O*: asymptotic convergence as input grows toward **infinity**

Reality: finite, often of quite small a scale

Big-*O*: assumes uniform computation cost

Reality: varied, with interferences, steps, and saturation points

Sure, but what's the harm?

- Gravitating toward algorithms that are theoretically good but practically bad
- Thinking there is “one true winner” for all scenarios

What to do (instead)?

Decide the scale range of the problem: $S: [s_{min}, s_{max}]$

Identify dominant operation(s): p , and their cost $C(p, s)$

Cost is a function of *both* operation p and scale s

Total cost model

Often a piecewise function

Dependent on multiple operations

What you should actually do

Start with just about any seemingly reasonable choice

Calculate cost leader(s) by tying resource utilization telemetry to “unit of work”, e.g. *rps*

Identify nature of bottleneck(s)—lower algorithm complexity matters only if compute is the bottleneck!